

Monitoreo y Observabilidad - Día 2

pending

40 min

Learning Objectives

- 1 Entender importancia del monitoreo en producción
- 2 Aprender métricas clave para pipelines
- 3 Comprender herramientas de observabilidad
- 4 Conocer estrategias de alerting efectivas

Theory

Practice

Evidence

Quiz

Practical exercise to apply the concepts learned.

Ejercicio: Configurar monitoreo completo para pipeline

Configurar métricas en DAG:

```
from airflow import DAG
from airflow.operators.python import PythonOperator
from airflow.metrics import Stats
from datetime import datetime, timedelta
import logging

logger = logging.getLogger(__name__)

def extract_with_metrics(**context):
    """Extracción con métricas detalladas"""
    try:
        # Simular extracción
        records = 1000
        duration = 45.2 # segundos

        # Registrar métricas
        Stats.gauge('pipeline.extract.records', records)
        Stats.timer('pipeline.extract.duration', duration)

        # Log estructurado
        logger.info(f"Extract completed: {records} records in {duration}s",
                   extra={'records': records, 'duration': duration})

        return {'records': records, 'duration': duration}

    except Exception as e:
        Stats.incr('pipeline.extract.errors')
        logger.error(f"Extract failed: {e}")
        raise e

def transform_with_metrics(**context):
    """Transformación con tracking de calidad"""
    ti = context['task_instance']
    input_data = ti.xcom_pull(task_ids='extract')

    try:
        # Simular transformación
        output_records = input_data['records'] * 0.95 # 5% pérdida
        quality_score = 0.98

        # Métricas de calidad
        Stats.gauge('pipeline.transform.quality', quality_score)
        Stats.gauge('pipeline.transform.output_records', output_records)

        # Alertas basadas en calidad
        if quality_score < 0.9:
            logger.warning(f"Low quality score: {quality_score}")

        return {'output_records': output_records, 'quality': quality_score}

    except Exception as e:
        Stats.incr('pipeline.transform.errors')
        raise e

def load_with_metrics(**context):
    """Carga con métricas de performance"""
    ti = context['task_instance']
    transform_data = ti.xcom_pull(task_ids='transform')

    try:
        # Simular carga
        records_loaded = transform_data['output_records']

        # Registro de métricas
        Stats.gauge('pipeline.load.records', records_loaded)
        Stats.gauge('pipeline.load.duration', duration)

        return {'records_loaded': records_loaded}

    except Exception as e:
        Stats.incr('pipeline.load.errors')
        raise e
```

▼

[↗ Sign Out]



```

load_time = 12.5

Stats.gauge('pipeline.load.records_loaded', records_loaded)
Stats.timer('pipeline.load.duration', load_time)

# Calcular throughput
throughput = records_loaded / load_time # records/second
Stats.gauge('pipeline.load.throughput', throughput)

return {'loaded': records_loaded, 'throughput': throughput}

except Exception as e:
    Stats.incr('pipeline.load.errors')
    raise e

# Callbacks de monitoreo
def dag_success_callback(context):
    """Callback cuando DAG completa exitosamente"""
    dag_run = context['dag_run']
    duration = (dag_run.end_date - dag_run.start_date).total_seconds()

    Stats.timer('pipeline.dag.duration', duration)
    Stats.incr('pipeline.dag.success')

    logger.info(f"DAG completed successfully in {duration}s")

def dag_failure_callback(context):
    """Callback cuando DAG falla"""
    dag_run = context['dag_run']
    Stats.incr('pipeline.dag.failure')

    # Enviar alerta crítica
    logger.error(f"DAG failed: {context['dag'].dag_id}")

def sla_miss_callback(dag, task_list, blocking_task_list, slas, blocking_tis):
    """Callback cuando se viola SLA"""
    Stats.incr('pipeline.sla.violations')
    logger.warning(f"SLA missed for DAG: {dag.dag_id}")

# Configurar DAG con monitoreo completo
dag = DAG(
    'pipeline_monitoring_demo',
    description='Pipeline con métricas y monitoreo completos',
    schedule_interval='@daily',
    start_date=datetime(2024, 1, 1),
    catchup=False,
    default_args={
        'retries': 2,
        'retry_delay': timedelta(minutes=5),
        'execution_timeout': timedelta(hours=2),
        'sla': timedelta(hours=1) # SLA: 1 hora máximo
    },
    # Callbacks de monitoreo
    on_success_callback=dag_success_callback,
    on_failure_callback=dag_failure_callback,
    sla_miss_callback=sla_miss_callback,
    tags=['monitoring', 'production']
)

# Tareas del pipeline
extract_task = PythonOperator(
    task_id='extract',
    python_callable=extract_with_metrics,
    provide_context=True,
    dag=dag
)

transform_task = PythonOperator(
    task_id='transform',
    python_callable=transform_with_metrics,
    provide_context=True,
    dag=dag
)

load_task = PythonOperator(
    task_id='load',
    python_callable=load_with_metrics,
    provide_context=True,
    dag=dag
)

# Dependencias
extract_task >> transform_task >> load_task

```

Configurar dashboard en Grafana:

```

// Panel de métricas principales
{
  "title": "Pipeline Health Dashboard",
  "panels": [
    {
      "title": "DAG Success Rate",
      "type": "stat",

```

```

    "targets": [
      "expr": "rate(pipeline_dag_success[1h]) / (rate(pipeline_dag_success[1h]) + rate(pipeline_dag_failure[1h])) * 100",
      "legendFormat": "Success Rate %"
    ]
  },
  {
    "title": "Average Processing Time",
    "type": "graph",
    "targets": [
      "expr": "rate(pipeline_extract_duration_sum[5m]) / rate(pipeline_extract_duration_count[5m])",
      "legendFormat": "Extract Time (s)"
    ]
  },
  {
    "title": "Error Rate by Component",
    "type": "bargauge",
    "targets": [
      {"expr": "rate(pipeline_extract_errors[1h])", "legendFormat": "Extract"},
      {"expr": "rate(pipeline_transform_errors[1h])", "legendFormat": "Transform"},
      {"expr": "rate(pipeline_load_errors[1h])", "legendFormat": "Load"}
    ]
  }
]
}

```

Configurar alertas:

```

# Alert rules para Prometheus
groups:
- name: pipeline_alerts
  rules:

  - alert: PipelineHighErrorRate
    expr: rate(pipeline_dag_failure[5m]) / (rate(pipeline_dag_success[5m]) + rate(pipeline_dag_failure[5m])) > 0.1
    for: 5m
    labels:
      severity: warning
    annotations:
      summary: "High error rate in pipeline"
      description: "Pipeline error rate > 10% for 5 minutes"

  - alert: PipelineSLAViolation
    expr: increase(pipeline_sla_violations[1h]) > 0
    labels:
      severity: critical
    annotations:
      summary: "SLA violation detected"
      description: "Pipeline failed to meet SLA requirements"

  - alert: PipelineDataFreshness
    expr: pipeline_data_freshness_hours > 24
    labels:
      severity: warning
    annotations:
      summary: "Data freshness issue"
      description: "Pipeline data is more than 24 hours old"

```

Verificación: ¿Qué métricas son más importantes para monitorear en un pipeline de datos vs una aplicación web? ¿Cómo decidir cuándo escalar de warning a critical en las alertas?

Requerimientos:

Apache Airflow con métricas habilitadas
 Prometheus/Grafana para monitoreo
 Configuración de logging estructurado

